

kmap manual

Introduction:

Kmap and kvector were designed as replacements for `std::map` and `std::vector` for handling large sized data structures. From testing, `std::map` and `std::vector` were found to be inefficient at memory management when storing large amounts of data. After further research the issue was found to be due to `std::allocator` which is a pool based allocator. Pool allocators store more memory than is actually needed by the application which allows memory allocation for the application to be faster. Unfortunately for large data structures, the memory bloat can fill up all of your ram. Kmap and kvector were designed to store large data structures as efficiently as possible.

Kvector also solved a key issue that `std::vector` does not do a very good job of solving. When `std::vector` is cleared, the allocated memory for the vector is not cleared. The only way to clear the memory for `std::vector` is to swap an empty vector with the full vector like so:

```
std::vector v1;  
v1.resize(200);  
{  
    std::vector v2;  
    v1.swap(v2);  
}
```

Kvector allows the allocated memory to be cleared with one command.

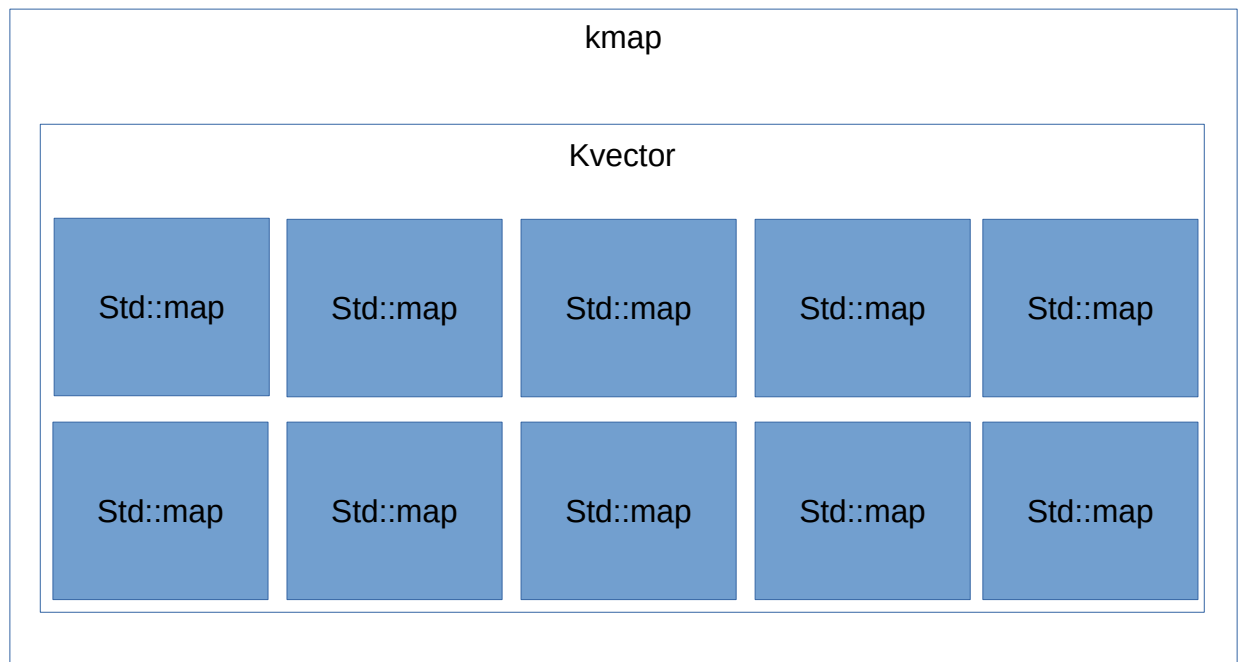


Figure 1: Diagram of the kmap data structure showing kvector storing `std::map`.

Kmap combines kvector with `std::map` (see Figure 1). Each `std::map` structure is designed to contain approximately 64 keys. This creates two advantages; one, the data structure is relatively flat, and two, `std::allocator` works extremely efficiently for `std::map`. Additionally, kmap uses a hash table to insert the key-value pair in the appropriate `std::map`. The insert operation is assumed to be $O(\log(64))$, since the index in kvector is assigned using a hashing function and there are at most 64 keys in each `std::map`. The read operation is amortized $O(1)$ because on average Y steps are needed to find the key-value pairs where Y is independent of the N key-value pairs stored in kvector. This occurs because the key-value pairs are hashed and `std::map` contains on average 64 key-value pairs under a well distributed hashing.

Kmap uses an iterator called `kmap_iterator`. `Kmap_iterator` stores the index of kvector and an iterator pointing to a key-value pair stored in `std::map`. Unlike iterators for `std::map` and `std::vector`, this iterator is not designed to be used with methods contained in the algorithms header, part of the c++ standard library.

kmap_iterator Methods:

❖ `bool operator != (uint64_t)`: tests if `kmap_iterator` has reached the end of kvector.

kmap Methods:

- ❖ `kmap()`: constructor that creates the default sized kmap
- ❖ `kmap(uint64_t)`: constructor that creates a specifically sized kmap
- ❖ `kmap(const kmap<K,V>&)`: copy constructor
- ❖ `kmap<K,V>& operator=(const kmap<K,V>&)`: assignment operator
- ❖ `void clear()`: replaces old kmap with an empty default sized kmap
- ❖ `void clean()`: clears all key-value pairs from kmap
- ❖ `void insert(const K&, const V&)`: inserts key-value pair into kmap
- ❖ `V& operator[](const K&)`: if key exists inserts value into kmap otherwise returns the value assigned to the key
- ❖ `void resize(uint64_t)`: sizes kmap to handle the number in the resize command. If the number of entries the current kmap can handle is greater than the value given in the command nothing happens. Otherwise kmap is resized and rehashed.
- ❖ `void swap (kmap &)`: swap between two kmaps
- ❖ `kmap_iterator<K,V> findkey(const K&)`: returns a `kmap_iterator` containing the hash value of the key and a `std::map` iterator pointing to the key in `std::map`.
- ❖ `void removekey(const K&)`: removes the key from kmap. If the key does not exist, no changes will occur to the kmap.
- ❖ `const K& getkey(const kmap_iterator<K,V> &)`: gets the key from kmap corresponding to the current value of `kmap_iterator`. This will throw an error if the iterator is at the end of `std::map`. This method is designed to be combined with kmap's iterating methods or `findkey`.
- ❖ `V& getvalue(const kmap_iterator<K,V> &)`: gets the value from kmap corresponding to the current value of `kmap_iterator`. This will throw an error if the iterator is at the end of `std::map`. This method is designed to be combined with kmap's iterating methods or `findkey`.
- ❖ `void setvalue(const kmap_iterator<K,V> &, const V&)`: sets the value from kmap corresponding to the current value `kmap_iterator`. This will throw an error if the iterator

is at the end of `std::map`. This method is designed to be combined with `kmap`'s iterating methods or `findkey`.

- ❖ `kmap_iterator<K,V> start()`: Finds the first `std::map` in `kvector` that is not empty and sets `kmap_iterator` to the beginning of `std::map`. `Kmap_iterator` is also assigned the index of the first non-empty `std::map`. If none of the `std::maps` are filled, this method sets the index in `kmap_iterator` to the value returned by the method `end`.
- ❖ `uint64_t end()`: returns the position after the end of `kvector` in `kmap`. The returned value is also the capacity of the `kvector` in `kmap`.
- ❖ `void next(kmap_iterator<K,V> &)`: Sets `kmap_iterator` to the next value in `std::map`. If `std::map`'s next value is the end, goto the next filled `std::map` and set `kmap_iterator` to the beginning of the `std::map`.
- ❖ `bool empty()`: returns true if `kmap` is empty. Otherwise, return false.

kvector Methods:

- ❖ `kvector()`: default constructor which sets `kvector` to the default size
- ❖ `kvector(uint64_t)`: constructor which sets the size of `kvector` to the specified size
- ❖ `kvector(uint64_t, T)`: constructor which sets the size of `kvector` the specified size and fills the `kvector` with the specified value.
- ❖ `kvector(const kvector &)`: copy constructor
- ❖ `void resize(uint64_t)`: resize `kvector` to be the specified size. Note: this can shrink the `kvector` which can cause loss of information but will not cause a memory leak.
- ❖ `void clear()`: deletes `kvector` and creates a new `kvector` with size and capacity of 0.
- ❖ `T & operator [] (uint64_t)`: returns the value stored at the `kvector`'s index.
- ❖ `const T & operator [] (uint64_t) const`: returns an unmodifiable value stored at the `kvector`'s index.
- ❖ `void push_back(T)`: add a value to the end of the `kvector` and increase the size of the `kvector` by one.
- ❖ `void pop_back()`: removes the value at the end of the `kvector` and decrease the size of the `kvector` by one.
- ❖ `uint64_t getsize() const`: returns the size of `kvector`.
- ❖ `uint64_t getcapacity() const`: returns the capacity of `kvector`.
- ❖ `bool empty() const`: returns true if the `kvector` is empty. Otherwise, this method returns false.
- ❖ `void clean()`: deletes the previous `kvector` and creates a new `kvector` with the same capacity as the old `kvector`.
- ❖ `void swap(kvector<T> &)`: swap between two `kvectors`
- ❖ `kvector<T> & operator= (const kvector<T> &)`: assignment operator
- ❖ `T & back()`: returns the value at the end of `kvector`.
- ❖ `const T & back() const`: returns an unmodifiable value at the end of `kvector`.
- ❖ `T & front()`: returns a value at the front of `kvector`.
- ❖ `const T & front() const`: returns an unmodifiable value at the front of `kvector`.